Qdrant is an open source vector database written in Rust.
It stores embeddings together with a JSON payload and supports filtering.
Collections use cosine, dot product or euclidean distance.

PDF parsing extracts the raw text of every page of a document.
The pypdf library reads PDF files in pure Python without native dependencies.

Chunking splits long documents into smaller overlapping pieces.
Each chunk is converted into a vector by an embedding model.
Retrieval returns the top k chunks closest to the question vector.

The weather in the mountains was cold and snowy during the winter holiday.
We cooked pasta with tomato sauce and drank hot chocolate by the fire.